


Lab 1 Back-propagation

313551055 柯柏旭


1. Introduction

在本報告中，我們將探討使用反向傳播算法訓練神經網絡的過程與結果。在 section [4.Discussion](#) 和 [5.Extra](#)，我們將利用不同的學習率、激活函數、和優化器等等來訓練，目的是評估這些因素對模型性能的影響。通過詳細的測試和討論，我們將展示這些方法如何影響網絡在處理**Linear數據**和**XOR數據**時的學習和預測能力，以下為Linear數據和XOR數據的呈現方式:

- Linear data



- XOR data



對於 section `2.Experiment setups` 和 `3.Result of testing` 皆由 base code 的 `main.py` 來陳述。

至於 section `4.Discussion` 和 `5.Extra` 則會另外寫成 `diff_` 開頭的 codes。

2. Experiment setups

Hardware overview

```
Model Name: MacBook Pro
Model Identifier: Mac15,6
Model Number: MRX33TA/A
Chip: Apple M3 Pro
Total Number of Cores: 11 (5 performance and 6 efficiency)
Memory: 18 GB
```

Python version

```
Python 3.9.19
```

A. Sigmoid functions

```
def sigmoid(x):
    return 1.0 / (1.0 + np.exp(-x))

def derivative_sigmoid(x):
    return np.multiply(x, 1.0 - x)
```

(refer to lab1 document)

B. Neural network

```
def forward_propagation(X, weights):
    Z1 = np.dot(X, weights['W1']) + weights['b1']
    A1 = sigmoid(Z1)
    Z2 = np.dot(A1, weights['W2']) + weights['b2']
    A2 = sigmoid(Z2)
    Z3 = np.dot(A2, weights['W3']) + weights['b3']
    A3 = sigmoid(Z3)
    return Z1, A1, Z2, A2, Z3, A3
```

Each hidden layers is $y = wx + b$ and *sigmoid* function

C. Backpropagation

```
def back_propagation(X, y, weights, Z1, A1, Z2, A2, Z3, A3,
                    learning_rate):
    m = y.shape[0] # number of samples

    # compute gradients
    dZ3 = (A3 - y) * 2
    dW3 = np.dot(A2.T, dZ3) / m
    db3 = np.sum(dZ3, axis=0, keepdims=True) / m # 沿著 row 求和，並保持
    維度

    dA2 = np.dot(dZ3, weights['W3'].T)
    dZ2 = dA2 * derivative_sigmoid(A2)
    dW2 = np.dot(A1.T, dZ2) / m
    db2 = np.sum(dZ2, axis=0, keepdims=True) / m
```

```

dA1 = np.dot(dZ2, weights['W2'].T)
dZ1 = dA1 * derivative_sigmoid(A1)
dW1 = np.dot(X.T, dZ1) / m
db1 = np.sum(dZ1, axis=0, keepdims=True) / m

# update weights
weights['W3'] -= learning_rate * dW3
weights['b3'] -= learning_rate * db3
weights['W2'] -= learning_rate * dW2
weights['b2'] -= learning_rate * db2
weights['W1'] -= learning_rate * dW1
weights['b1'] -= learning_rate * db1

```

拿 $Z3$, $A3$, $W3$ 為例，推導如下：

首先定義 $dZ3$ 如下

$$dZ3 = \frac{\partial L}{\partial Z3}, dZ3 = 2 * (A3 - y)$$

接著我們需要計算損失函數 (L) 對 權重($W3$) 的梯度 ($dW3 = \frac{\partial L}{\partial W3}$)。

由於 ($Z3 = A2 \cdot W3 + b3$)，我們可以將 (L) 對 ($W3$) 的偏微分寫成：

$$\frac{\partial L}{\partial W3} = \frac{\partial L}{\partial Z3} \cdot \frac{\partial Z3}{\partial W3}$$

然後運用 chain rule，

首先計算 ($\frac{\partial Z3}{\partial W3}$)：

$$Z3 = A2 \cdot W3 + b3$$

$$\frac{\partial Z3}{\partial W3} = A2$$

將上面的結果代入鏈式法則：

$$\frac{\partial L}{\partial W3} = \frac{\partial L}{\partial Z3} \cdot A2$$

由於 ($dZ3 = \frac{\partial L}{\partial Z3}$)，我們可以得到：

$$dW3 = dZ3 \cdot A2$$

D. Weight Initialization

```
def initialize_weights(input_size, hidden_size1, hidden_size2,
                        output_size):
    weights = {
        'W1': np.random.randn(input_size, hidden_size1), # shape:
(input_size, hidden_size1)
        'b1': np.zeros((1, hidden_size1)), # shape: (1, hidden_size1)
        'W2': np.random.randn(hidden_size1, hidden_size2),
        'b2': np.zeros((1, hidden_size2)),
        'W3': np.random.randn(hidden_size2, output_size),
        'b3': np.zeros((1, output_size))
    }
    return weights
```

*weights*大多會初始化在 $[-3, 3]$ 之間

*bias*皆初始化為0

E. Loss function

```
# cost(loss) function: MSE
def compute_loss(y_true, y_pred):
    return np.mean((y_true - y_pred) ** 2)
```

loss (cost) function 為傳統 L2 norm

F. Train

```
def train(X, y, input_size, hidden_size1, hidden_size2, output_size,
learning_rate, epochs):
    weights = initialize_weights(input_size, hidden_size1,
hidden_size2, output_size)

    for epoch in range(epochs):
        Z1, A1, Z2, A2, Z3, A3 = forward_propagation(X, weights)
        loss = compute_loss(y, A3)
        back_propagation(X, y, weights, Z1, A1, Z2, A2, Z3, A3,
learning_rate)


        if epoch % 5000 == 0:
            print(f'Epoch {epoch}, Loss: {loss}')

    return weights
```


3. Result of testing

A. Screenshot and comparison figure

- Linear data



- XOR data



B. Show the accuracy of your prediction

- Linear data (acc = 100%)

```

Iter1 | Ground truth: 0 | prediction: 0.00000 |
Iter2 | Ground truth: 1 | prediction: 1.00000 |
Iter3 | Ground truth: 0 | prediction: 0.00000 |
Iter4 | Ground truth: 1 | prediction: 1.00000 |
Iter5 | Ground truth: 0 | prediction: 0.00000 |
Iter6 | Ground truth: 0 | prediction: 0.00000 |
Iter7 | Ground truth: 0 | prediction: 0.00000 |
Iter8 | Ground truth: 0 | prediction: 0.00000 |
Iter9 | Ground truth: 1 | prediction: 1.00000 |
Iter10 | Ground truth: 1 | prediction: 1.00000 |
Iter11 | Ground truth: 1 | prediction: 1.00000 |
Iter12 | Ground truth: 1 | prediction: 1.00000 |

```


Iter13		Ground truth: 1		prediction: 1.00000	
Iter14		Ground truth: 1		prediction: 1.00000	
Iter15		Ground truth: 1		prediction: 1.00000	
Iter16		Ground truth: 1		prediction: 1.00000	
Iter17		Ground truth: 0		prediction: 0.00000	
Iter18		Ground truth: 1		prediction: 1.00000	
Iter19		Ground truth: 1		prediction: 1.00000	
Iter20		Ground truth: 1		prediction: 1.00000	
Iter21		Ground truth: 0		prediction: 0.00000	
Iter22		Ground truth: 1		prediction: 0.98980	
Iter23		Ground truth: 1		prediction: 1.00000	
Iter24		Ground truth: 0		prediction: 0.02334	
Iter25		Ground truth: 0		prediction: 0.00000	
Iter26		Ground truth: 0		prediction: 0.00000	
Iter27		Ground truth: 0		prediction: 0.00000	
Iter28		Ground truth: 0		prediction: 0.00000	
Iter29		Ground truth: 1		prediction: 1.00000	
Iter30		Ground truth: 1		prediction: 1.00000	
Iter31		Ground truth: 1		prediction: 1.00000	
Iter32		Ground truth: 0		prediction: 0.00000	
Iter33		Ground truth: 0		prediction: 0.00002	
Iter34		Ground truth: 1		prediction: 1.00000	
Iter35		Ground truth: 0		prediction: 0.00000	
Iter36		Ground truth: 0		prediction: 0.00000	
Iter37		Ground truth: 0		prediction: 0.00000	
Iter38		Ground truth: 0		prediction: 0.00000	
Iter39		Ground truth: 1		prediction: 0.99978	
Iter40		Ground truth: 0		prediction: 0.00000	
Iter41		Ground truth: 1		prediction: 1.00000	
Iter42		Ground truth: 1		prediction: 1.00000	
Iter43		Ground truth: 0		prediction: 0.00000	
Iter44		Ground truth: 0		prediction: 0.00000	
Iter45		Ground truth: 0		prediction: 0.00000	
Iter46		Ground truth: 1		prediction: 1.00000	
Iter47		Ground truth: 0		prediction: 0.00003	
Iter48		Ground truth: 0		prediction: 0.00000	
Iter49		Ground truth: 0		prediction: 0.00000	

Iter50		Ground truth: 0		prediction: 0.00006	
Iter51		Ground truth: 1		prediction: 1.00000	
Iter52		Ground truth: 0		prediction: 0.00000	
Iter53		Ground truth: 0		prediction: 0.00000	
Iter54		Ground truth: 1		prediction: 1.00000	
Iter55		Ground truth: 1		prediction: 1.00000	
Iter56		Ground truth: 1		prediction: 1.00000	
Iter57		Ground truth: 1		prediction: 1.00000	
Iter58		Ground truth: 1		prediction: 1.00000	
Iter59		Ground truth: 0		prediction: 0.00000	
Iter60		Ground truth: 1		prediction: 1.00000	
Iter61		Ground truth: 0		prediction: 0.00000	
Iter62		Ground truth: 0		prediction: 0.00000	
Iter63		Ground truth: 1		prediction: 1.00000	
Iter64		Ground truth: 0		prediction: 0.00000	
Iter65		Ground truth: 0		prediction: 0.00000	
Iter66		Ground truth: 1		prediction: 1.00000	
Iter67		Ground truth: 0		prediction: 0.00000	
Iter68		Ground truth: 0		prediction: 0.00000	
Iter69		Ground truth: 1		prediction: 0.99995	
Iter70		Ground truth: 1		prediction: 1.00000	
Iter71		Ground truth: 1		prediction: 1.00000	
Iter72		Ground truth: 1		prediction: 1.00000	
Iter73		Ground truth: 0		prediction: 0.00000	
Iter74		Ground truth: 0		prediction: 0.00000	
Iter75		Ground truth: 1		prediction: 1.00000	
Iter76		Ground truth: 1		prediction: 1.00000	
Iter77		Ground truth: 1		prediction: 1.00000	
Iter78		Ground truth: 1		prediction: 1.00000	
Iter79		Ground truth: 0		prediction: 0.00000	
Iter80		Ground truth: 1		prediction: 1.00000	
Iter81		Ground truth: 0		prediction: 0.00000	
Iter82		Ground truth: 1		prediction: 1.00000	
Iter83		Ground truth: 1		prediction: 0.98803	
Iter84		Ground truth: 1		prediction: 1.00000	
Iter85		Ground truth: 0		prediction: 0.00000	
Iter86		Ground truth: 1		prediction: 1.00000	

```
Iter87 | Ground truth: 1 | prediction: 1.00000 |
Iter88 | Ground truth: 1 | prediction: 1.00000 |
Iter89 | Ground truth: 0 | prediction: 0.00000 |
Iter90 | Ground truth: 0 | prediction: 0.00000 |
Iter91 | Ground truth: 1 | prediction: 1.00000 |
Iter92 | Ground truth: 0 | prediction: 0.00000 |
Iter93 | Ground truth: 0 | prediction: 0.00000 |
Iter94 | Ground truth: 1 | prediction: 1.00000 |
Iter95 | Ground truth: 0 | prediction: 0.00000 |
Iter96 | Ground truth: 1 | prediction: 1.00000 |
Iter97 | Ground truth: 1 | prediction: 1.00000 |
Iter98 | Ground truth: 0 | prediction: 0.00001 |
Iter99 | Ground truth: 1 | prediction: 1.00000 |
Iter100 | Ground truth: 1 | prediction: 1.00000 |
loss=0.0000079244 accuracy=100.00%
```

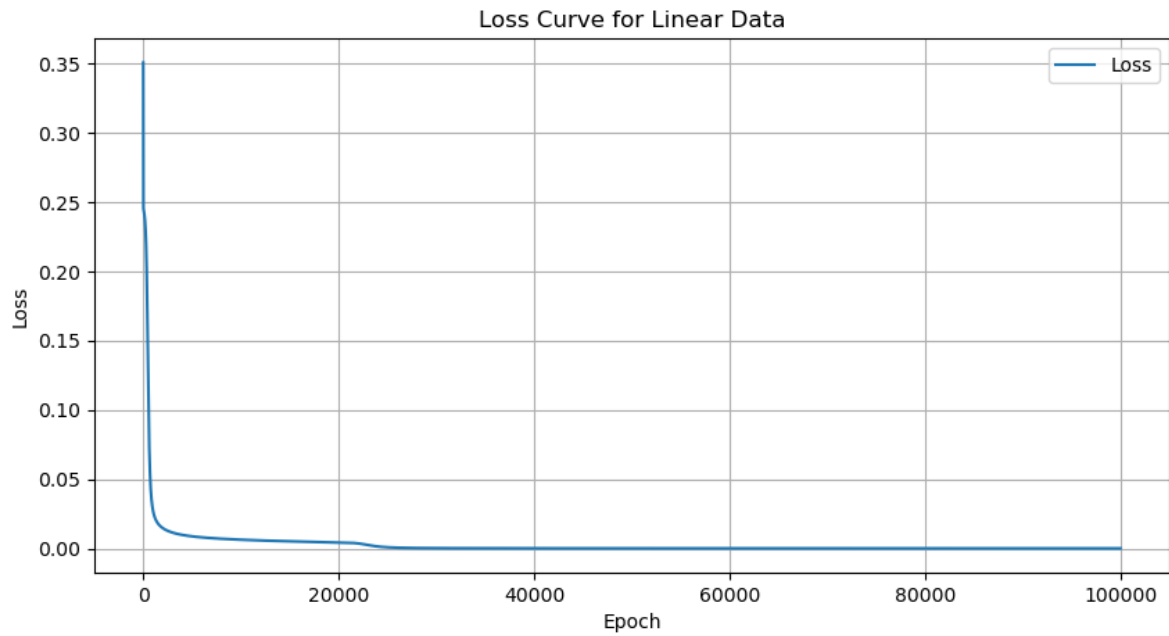
▪ XOR data (acc = 100%)

```
Iter1 | Ground truth: 0 | prediction: 0.00021 |
Iter2 | Ground truth: 1 | prediction: 0.99973 |
Iter3 | Ground truth: 0 | prediction: 0.00023 |
Iter4 | Ground truth: 1 | prediction: 0.99973 |
Iter5 | Ground truth: 0 | prediction: 0.00027 |
Iter6 | Ground truth: 1 | prediction: 0.99973 |
Iter7 | Ground truth: 0 | prediction: 0.00031 |
Iter8 | Ground truth: 1 | prediction: 0.99971 |
Iter9 | Ground truth: 0 | prediction: 0.00035 |
Iter10 | Ground truth: 1 | prediction: 0.99902 |
Iter11 | Ground truth: 0 | prediction: 0.00037 |
Iter12 | Ground truth: 0 | prediction: 0.00036 |
Iter13 | Ground truth: 1 | prediction: 0.99909 |
Iter14 | Ground truth: 0 | prediction: 0.00034 |
Iter15 | Ground truth: 1 | prediction: 0.99992 |
Iter16 | Ground truth: 0 | prediction: 0.00031 |
Iter17 | Ground truth: 1 | prediction: 0.99993 |
Iter18 | Ground truth: 0 | prediction: 0.00029 |
Iter19 | Ground truth: 1 | prediction: 0.99993 |
Iter20 | Ground truth: 0 | prediction: 0.00026 |
```


```
Iter21 | Ground truth: 1 | prediction: 0.99994 |  
loss=0.0000001487 accuracy=100.00%
```

C. Learning curve (loss, epoch curve)

■ Linear data



■ XOR data




D. Anything you want to present

4. Discussion

A. Try different learning rates


For linear data

- $lr = 1$, $acc = 100\%$




- $lr = 0.1$, $acc = 100\%$


Ground truth




Predict result (lr=0.1)



■ lr = 0.0001, acc = 52%




- Comparison figure




For XOR data

- $lr = 1$, acc = 100%




■ lr = 0.1, acc = 100%

Ground truth




Predict result (lr=0.1)




■ lr = 0.0001, acc = 52.38%


Ground truth



Predict result (lr=1e-05)



■ Comparison figure




Learning rate 設在 0.00001 太小了，原本以為是 epochs 不夠多無法讓他收斂，但觀察上面的比較圖後，看起來模型早在 10000 epoch 就沒有辦法降低 loss 了。



B. Try different numbers of hidden units

For linear data


- hidden units = (2, 2) , acc = 100%



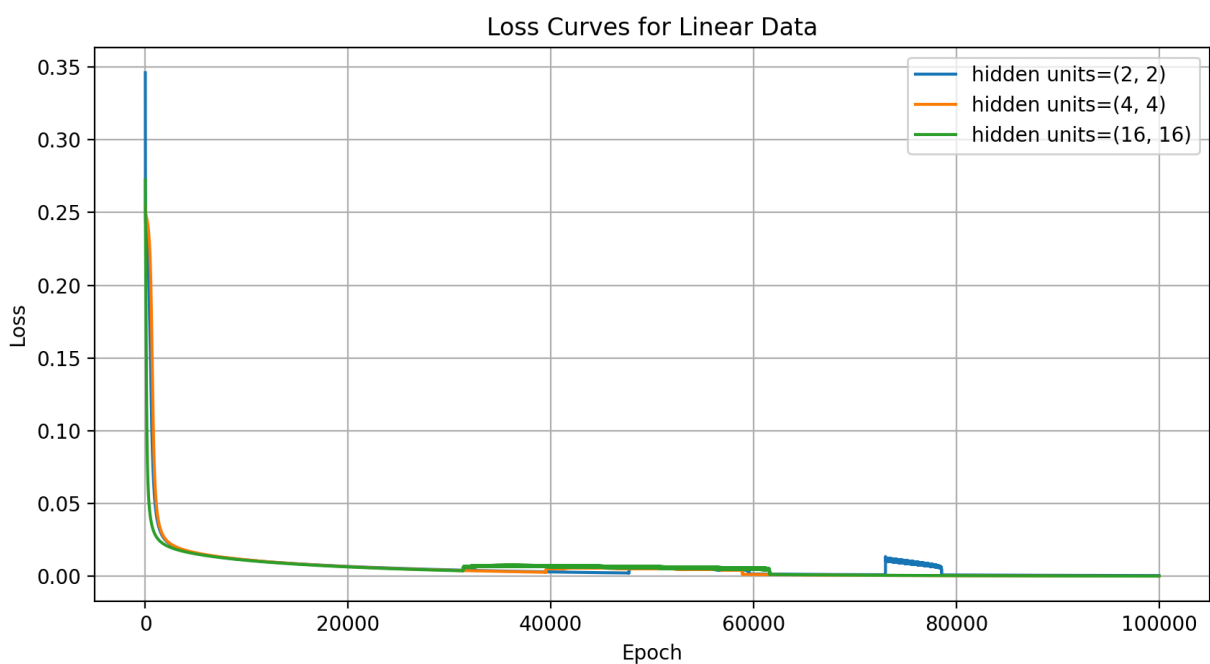
- hidden units = (4, 4) , acc = 100%



▪ hidden units = (16, 16) , acc = 100%




■ Comparison figure




For XOR data

- hidden units = (2, 2), acc = 100%

Ground truth




Predict result ((2, 2))




- hidden units = (4, 4), acc = 100%

Ground truth




Predict result ((4, 4))




- hidden units = (16, 16), acc = 100%

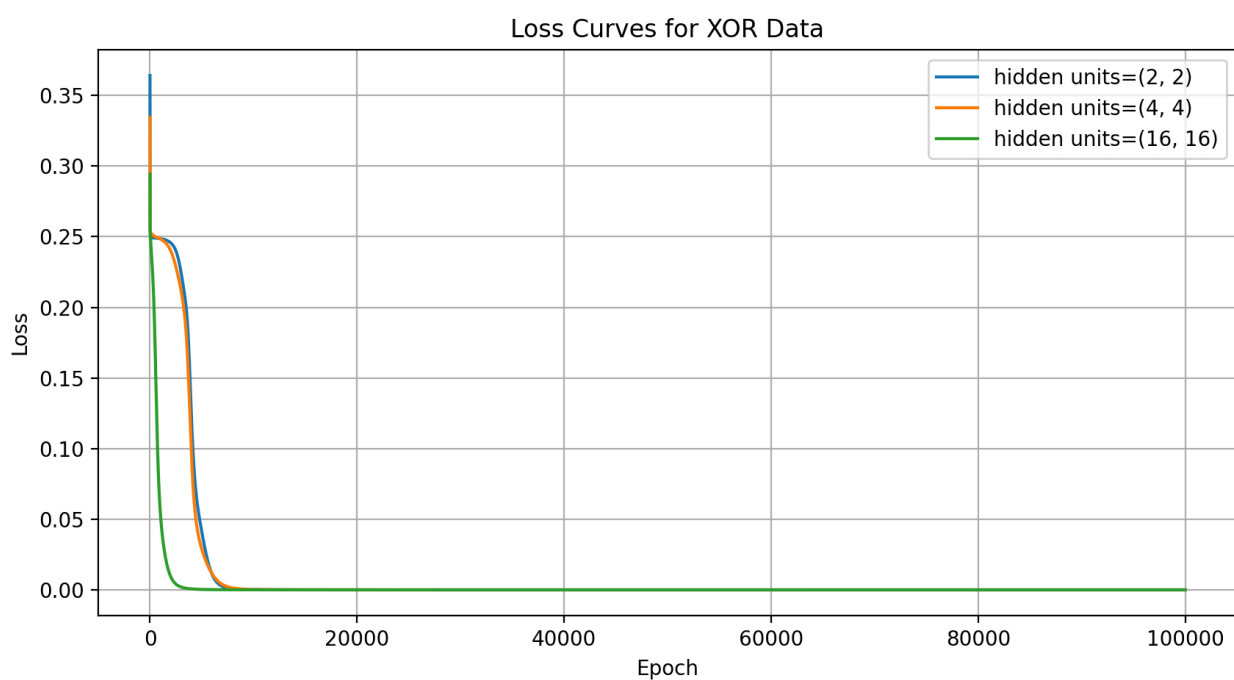
Ground truth



Predict result ((16, 16))



Comparison figure



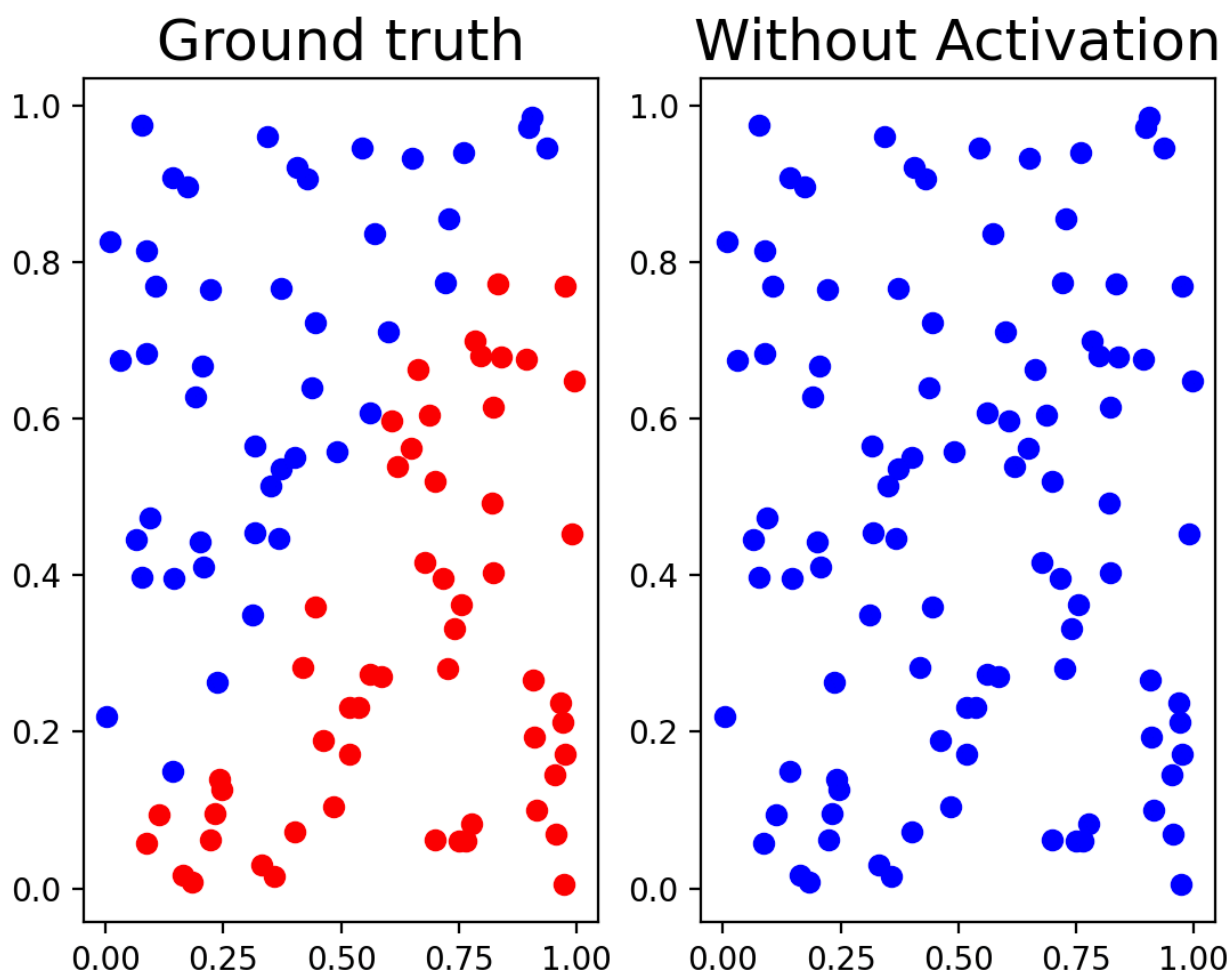
觀察下來，好像 hidden units 設在 (16, 16)，收斂速度會更快一些。

C. Try without activation functions

Without activation function (acc 還是能反映大致訓練狀況，但僅參考...)


- Linear data

少數情況會完全 train 不起來，acc = 55%




通常是都 train 得起來，acc = 98%，但還是找不到最佳解


Ground truth



Without Activation




Comparison figure




- XOR data

完全無法 train , acc = 52 %


Ground truth



Without Activation



Comparison figure



loss 會直接突破天際...

D. Anything you want to share

5. Extra

A. Implement different optimizers

Momentum


```
def momentum_update(weights, grads, velocity, learning_rate,
momentum=0.9):
    for key in weights.keys():
        velocity[key] = momentum * velocity[key] + learning_rate *
grads[key]
        weights[key] += velocity[key]
    return weights, velocity
```

`velocity[key]` 表示上一次更新的速度。

(如果更新方向和上次相反，這次更新速度會變慢，反之則變快。)


- Linear data (acc = 100%)

Ground truth Predict result (MOMENTUM)



■ XOR data (acc = 100%)

Ground truth Predict result (MOMENTUM)



Adam


```
def adam_update(weights, grads, m, v, t, learning_rate, beta1=0.9,
beta2=0.999, epsilon=1e-8):
    m_hat = {}
    v_hat = {}
    for key in weights.keys():
        m[key] = beta1 * m[key] + (1 - beta1) * grads[key]
        v[key] = beta2 * v[key] + (1 - beta2) * np.square(grads[key])
        m_hat[key] = m[key] / (1 - beta1 ** t)
        v_hat[key] = v[key] / (1 - beta2 ** t)
        weights[key] -= learning_rate * m_hat[key] /
(np.sqrt(v_hat[key]) + epsilon)
    return weights, m, v
```

其中：


- `weights`: 當前的權重參數。
- `grads`: 當前計算出的梯度。
- `m`: 一階矩估計（動量）。
- `v`: 二階矩估計（RMSProp）。
- `t`: 當前的時間步（迭代次數）。
- `learning_rate`: 學習率。
- `beta1`: 一階矩估計的衰減率，通常設為 0.9。
- `beta2`: 二階矩估計的衰減率，通常設為 0.999。
- `epsilon`: 防止除零的小數值，通常設為 $1e-8$ 。

■ Linear data (acc = 100%)

Ground truth




Predict result (ADAM)




■ XOR data (acc = 100%)

Ground truth




Predict result (ADAM)




Compare figure

- Linear data

Loss Curves for Linear Data



- XOR data



經過觀察可以發現，收斂速度為 **ADAM** > **MOMENTUM** > **SGD**

B. Implement different activation functions.

將 **sigmoid** 換成 **relu**


```
def relu(x):  
    return np.maximum(0, x)  
  
def derivative_relu(x):  
    return np.where(x > 0, 1, 0)
```

因為對於 ReLU，通常輸出值的範圍為 $[0, \infty)$ ，這邊的實驗將試著根據訓練數據的輸出值來動態調整閾值，即使用輸出值的中位數作為 threshold，但 acc 還是不能參考...


跟 sigmoid 比較起來，好像 relu 的收斂穩定度沒有它來得高。

- Linear data (acc = 98%)


Loss Curve for Linear Data






Ground truth



Predict result



■ XOR data (acc = 100%)



C. Implement convolutional layers.

Reference

- <https://hackmd.io/@allen108108/H1l4zqtp4> (Adagrad、RMSprop、Momentum and Adam – 特殊的學習率調整方式)
- <https://www.brilliantcode.net/1670/convolutional-neural-networks-4-backpropagation-in-kernels-of-cnns/> (卷積核的Back propagation)
- ChatGPT